



Lecture 7 – Real Time Systems (Sommerville Ch. 13)

Karl R. Wilcox
K.R.Wilcox@reading.ac.uk



Objectives

- To explain the concept of a real-time system and why these systems are usually implemented as concurrent processes
- To describe a design process for real-time systems
- To explain the role of a real-time executive
- To introduce generic architectures for monitoring and control and data acquisition systems



Real-time systems

- **Systems which monitor and control their environment**
- **Inevitably associated with hardware devices**
 - Sensors: Collect data from the system environment
 - Actuators: Change (in some way) the system's environment
- **Time is critical. Real-time systems MUST respond within specified times**



Definition

- A **real-time system** is a software system where the correct functioning of the system depends on the results produced by the system and the time at which these results are produced
- A ‘soft’ real-time system is a system whose operation is degraded if results are not produced according to the specified timing requirements
- A ‘hard’ real-time system is a system whose operation is incorrect if results are not produced according to the timing specification



Stimulus/Response Systems

- Given a stimulus, the system must produce a response within a specified time
- Periodic stimuli. Stimuli which occur at predictable time intervals
 - For example, a temperature sensor may be polled 10 times per second
- Aperiodic stimuli. Stimuli which occur at unpredictable times
 - For example, a system power failure may trigger an interrupt which must be processed by the system

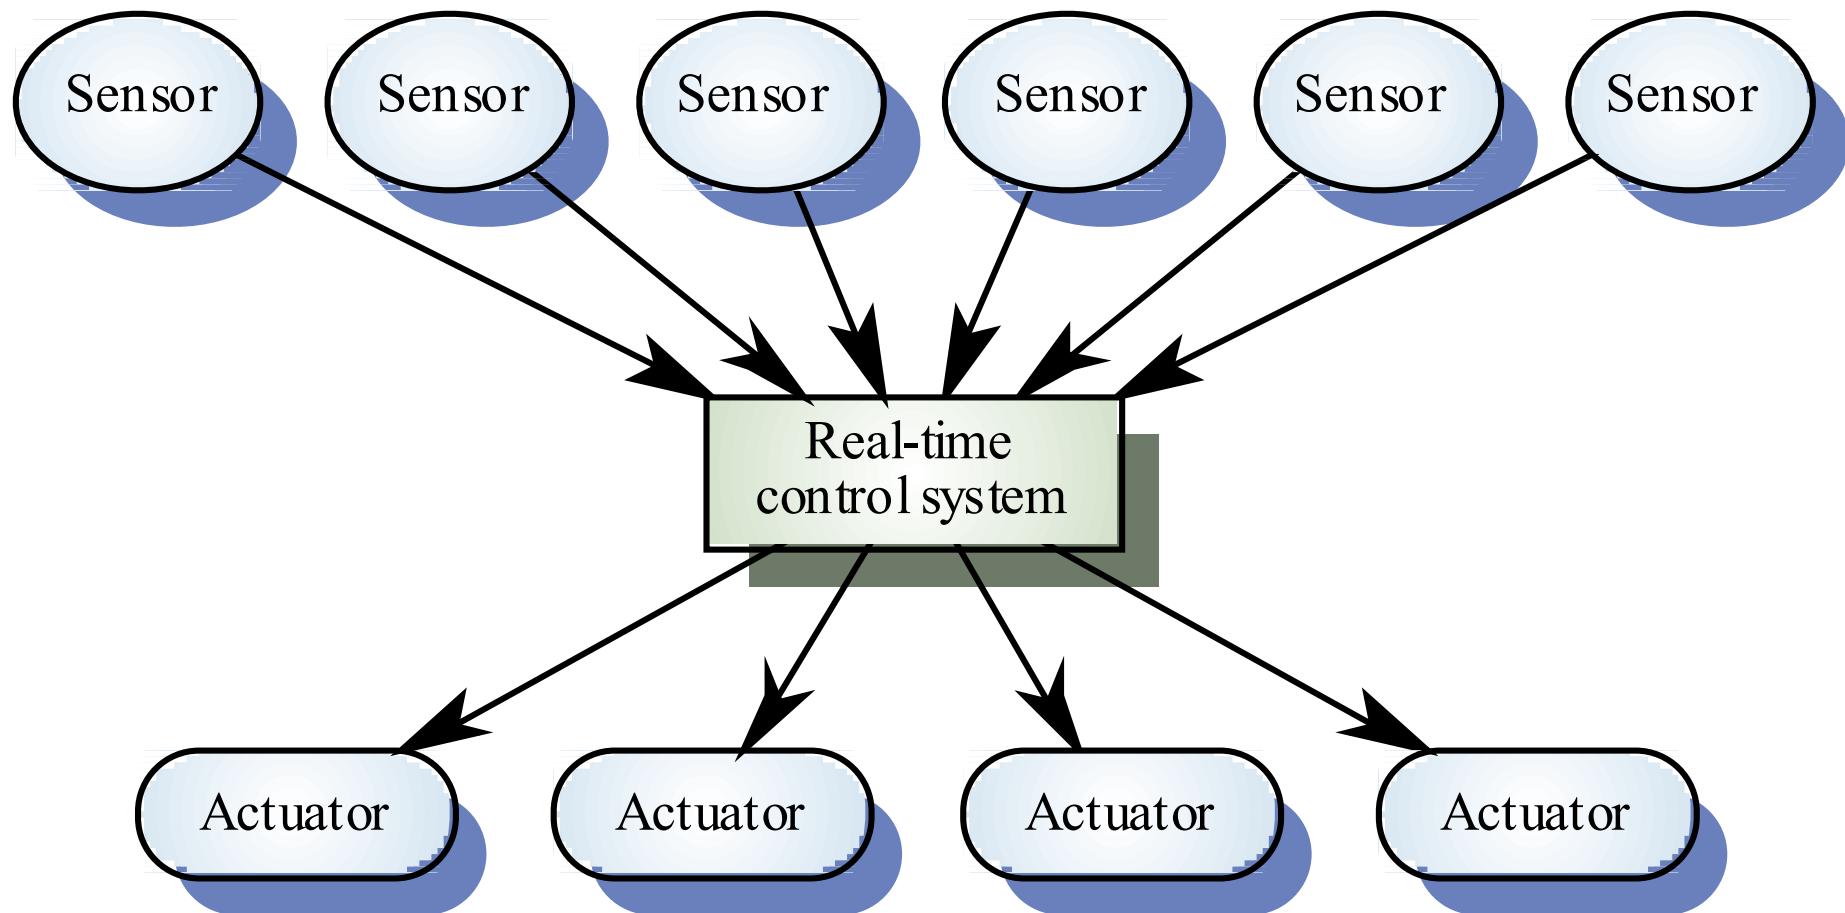


Architectural considerations

- Because of the need to respond to timing demands made by different stimuli/responses, the system architecture must allow for fast switching between stimulus handlers
- Timing demands of different stimuli are different so a simple sequential loop is not usually adequate
- Real-time systems are usually designed as cooperating processes with a real-time executive controlling these processes



A real-time system model



The University of Reading

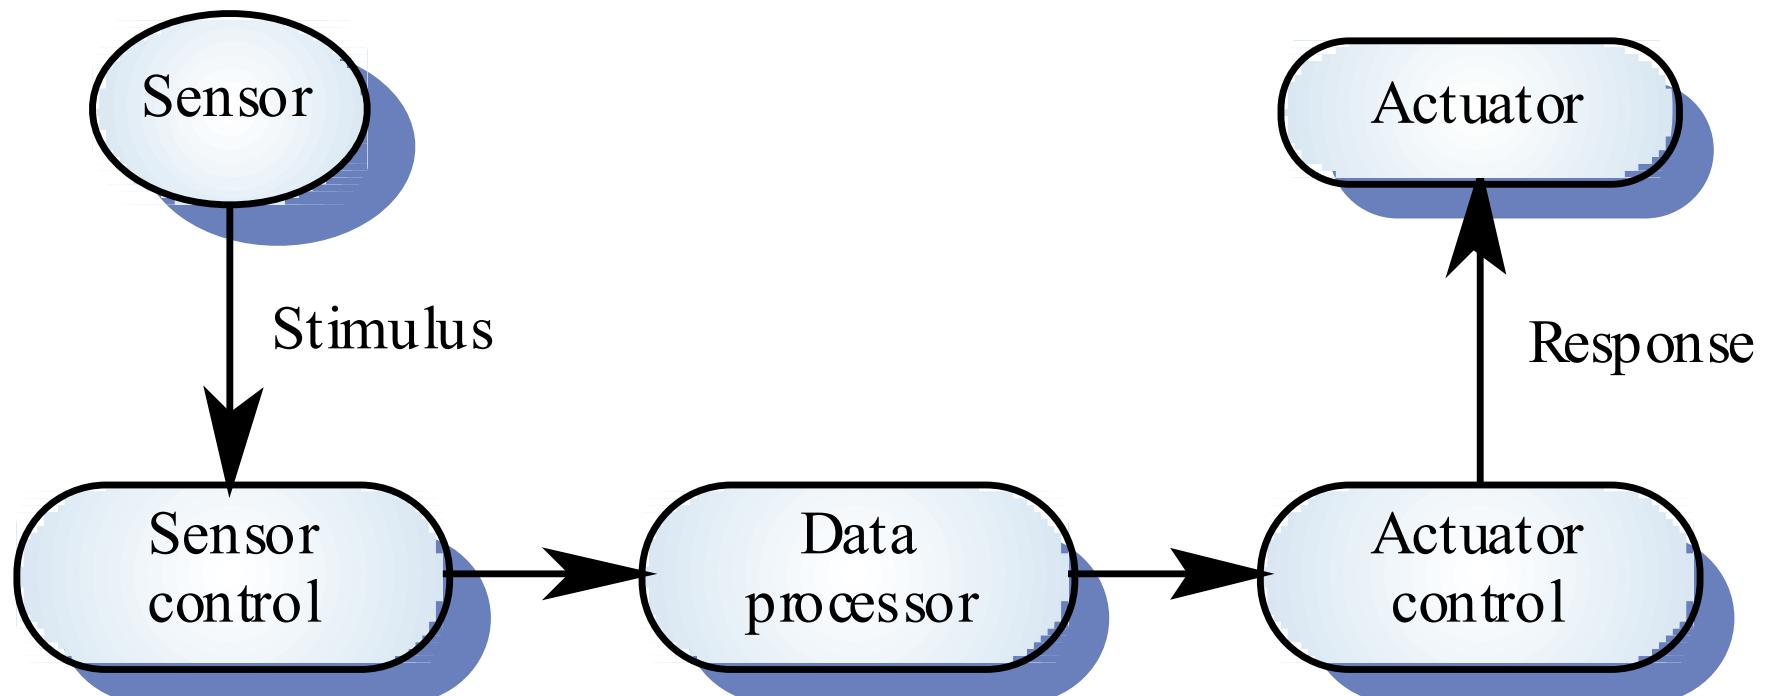


System elements

- **Sensors control processes**
 - Collect information from sensors. May buffer information collected in response to a sensor stimulus
- **Data processor**
 - Carries out processing of collected information and computes the system response
- **Actuator control**
 - Generates control signals for the actuator



Sensor/actuator processes



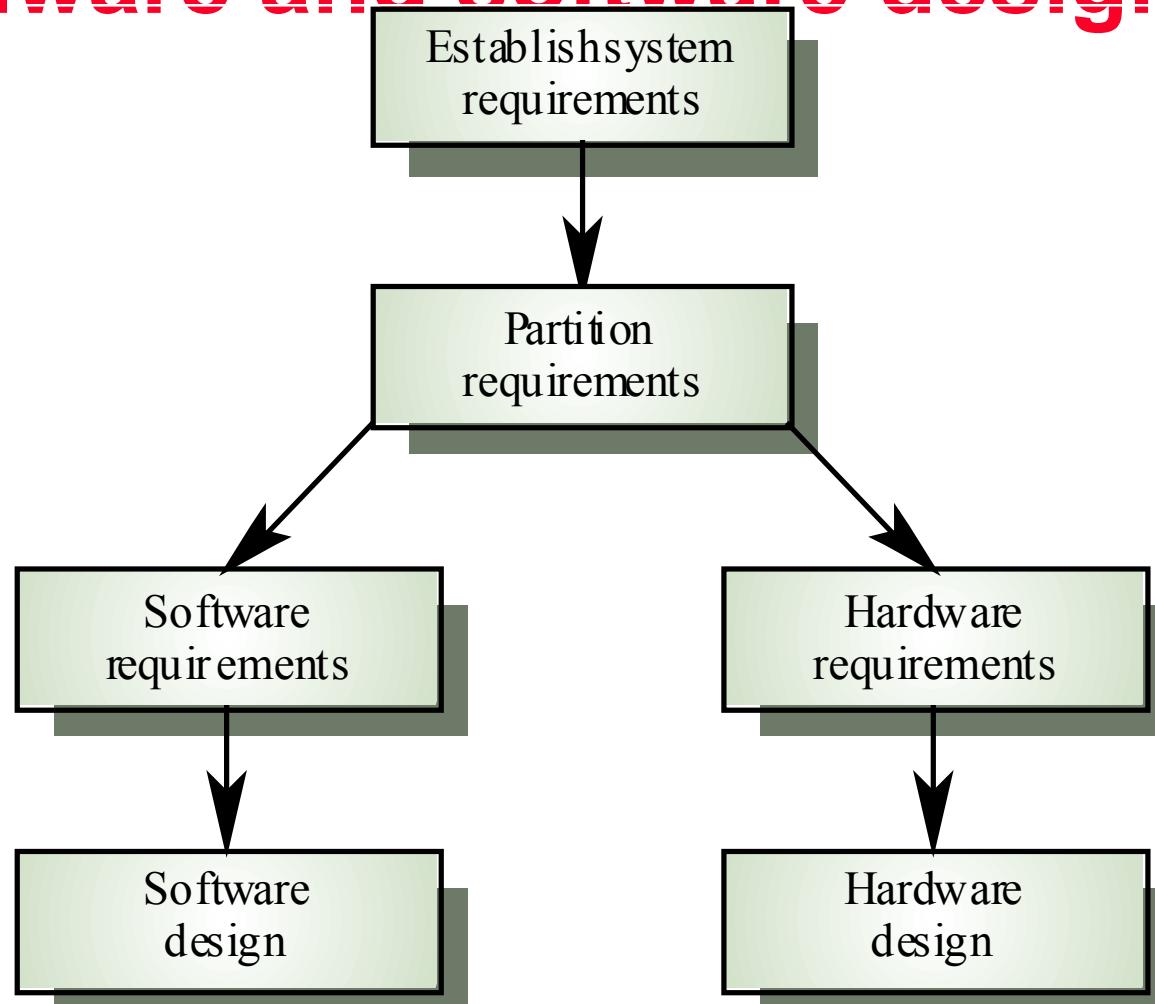


System design

- Design both the hardware and the software associated with system. Partition functions to either hardware or software
- Design decisions should be made on the basis on non-functional system requirements
- Hardware delivers better performance but potentially longer development and less scope for change



Hardware and software design





R-T systems design process

- Identify the stimuli to be processed and the required responses to these stimuli
- For each stimulus and response, identify the timing constraints
- Aggregate the stimulus and response processing into concurrent processes. A process may be associated with each class of stimulus and response



R-T systems design process

- Design algorithms to process each class of stimulus and response. These must meet the given timing requirements
- Design a scheduling system which will ensure that processes are started in time to meet their deadlines
- Integrate using a real-time executive or operating system



Timing constraints

- May require extensive simulation and experiment to ensure that these are met by the system
- May mean that certain design strategies such as object-oriented design cannot be used because of the additional overhead involved
- May mean that low-level programming language features have to be used for performance reasons

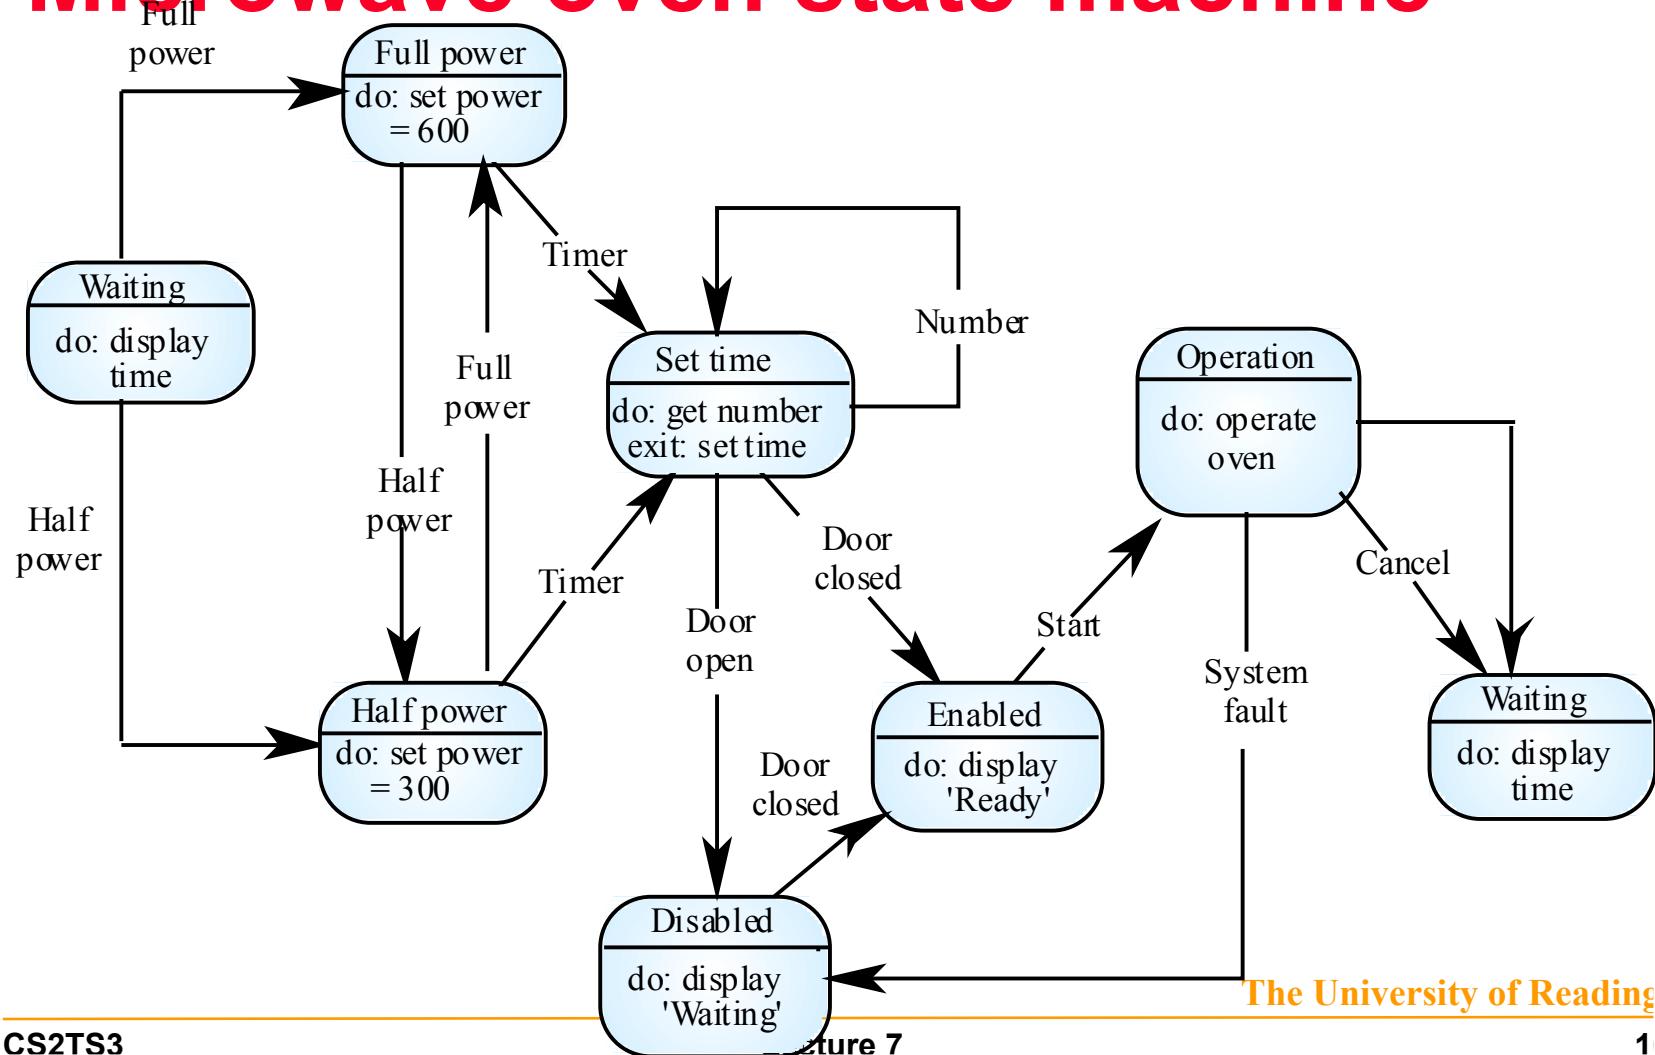


State machine modelling

- The effect of a stimulus in a real-time system may trigger a transition from one state to another.
- Finite state machines can be used for modelling real-time systems.
- However, FSM models lack structure. Even simple systems can have a complex model.
- The UML includes notations for defining state machine models
- See also Chapter 7.



Microwave oven state machine





Real-time programming

- Hard-real time systems may have to be programmed in assembly language to ensure that deadlines are met
- Languages such as C allow efficient programs to be written but do not have constructs to support concurrency or shared resource management
- Ada as a language designed to support real-time systems design so includes a general purpose concurrency mechanism



Java as a real-time language

- Java supports lightweight concurrency (threads and synchronized methods) and can be used for some soft real-time systems
- Java 2.0 is not suitable for hard RT programming or programming where precise control of timing is required
 - Not possible to specify thread execution time
 - Uncontrollable garbage collection
 - Not possible to discover queue sizes for shared resources
 - Variable virtual machine implementation
 - Not possible to do space or timing analysis



Real-time executives

- Real-time executives are specialised operating systems which manage the processes in the RTS
- Responsible for process management and resource (processor and memory) allocation
- May be based on a standard RTE kernel which is used unchanged or modified for a particular application
- Does not include facilities such as file management



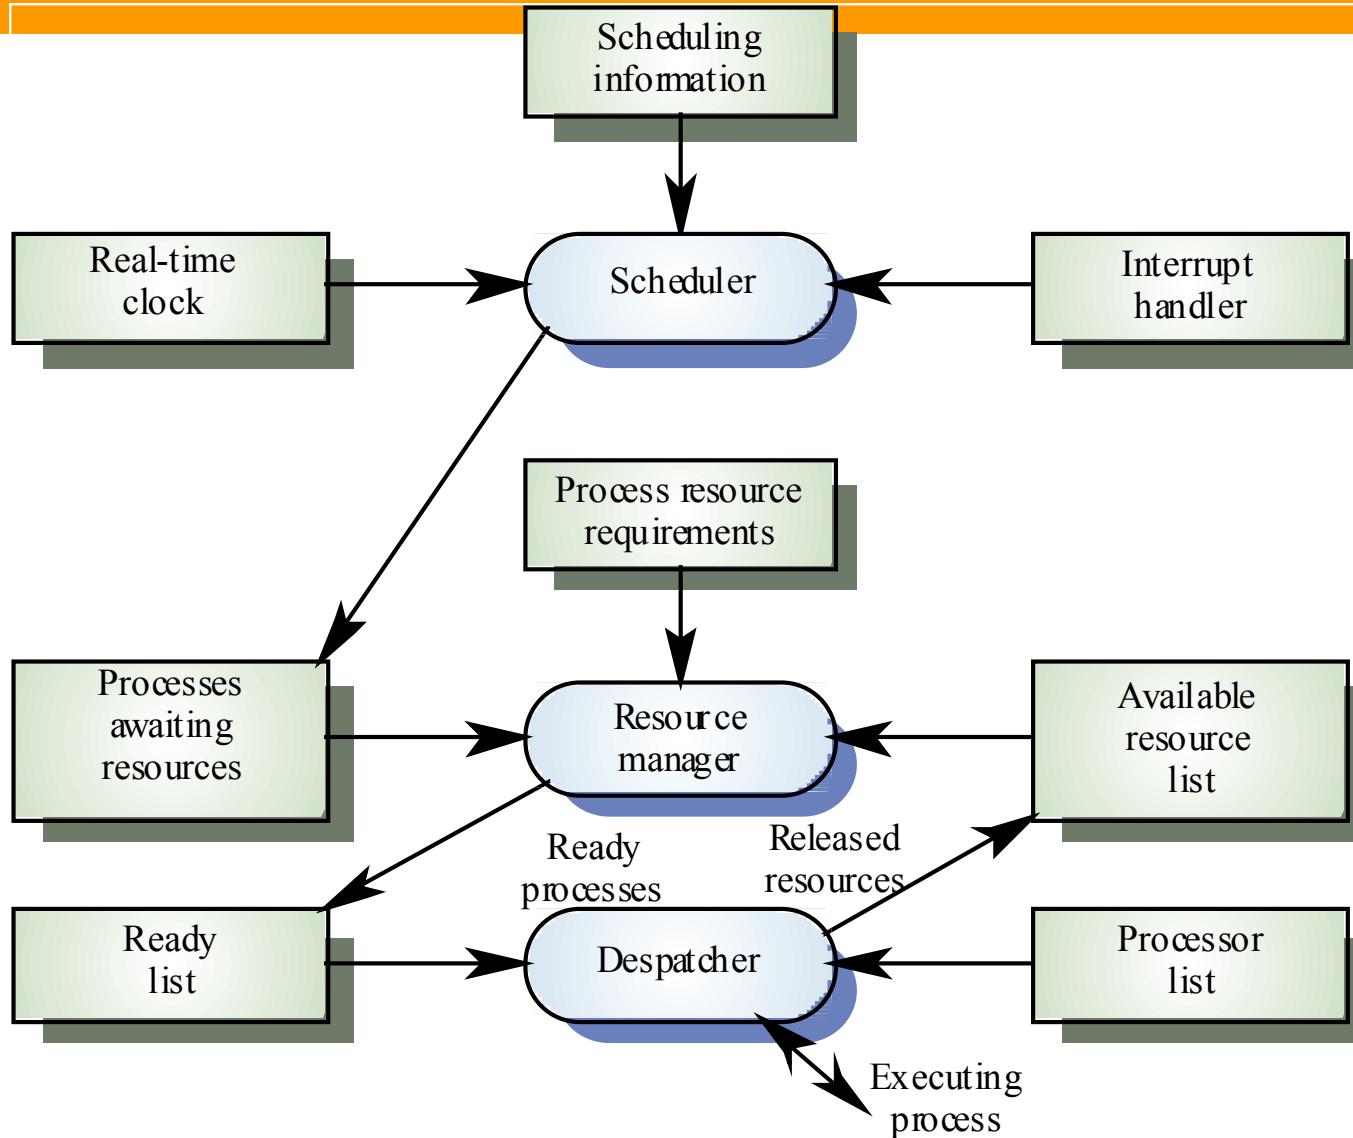
Executive components

- **Real-time clock**
 - Provides information for process scheduling.
- **Interrupt handler**
 - Manages aperiodic requests for service.
- **Scheduler**
 - Chooses the next process to be run.
- **Resource manager**
 - Allocates memory and processor resources.
- **Despatcher**
 - Starts process execution.



Non-stop system components

- **Configuration manager**
 - Responsible for the dynamic reconfiguration of the system software and hardware. Hardware modules may be replaced and software upgraded without stopping the systems
- **Fault manager**
 - Responsible for detecting software and hardware faults and taking appropriate actions (e.g. switching to backup disks) to ensure that the system continues in operation





Process priority

- The processing of some types of stimuli must sometimes take priority
- Interrupt level priority. Highest priority which is allocated to processes requiring a very fast response
- Clock level priority. Allocated to periodic processes
- Within these, further levels of priority may be assigned



Interrupt servicing

- Control is transferred automatically to a pre-determined memory location
- This location contains an instruction to jump to an interrupt service routine
- Further interrupts are disabled, the interrupt serviced and control returned to the interrupted process
- Interrupt service routines MUST be short, simple and fast



Periodic process servicing

- In most real-time systems, there will be several classes of periodic process, each with different periods (the time between executions), execution times and deadlines (the time by which processing must be completed)
- The real-time clock ticks periodically and each tick causes an interrupt which schedules the process manager for periodic processes
- The process manager selects a process which is ready for execution

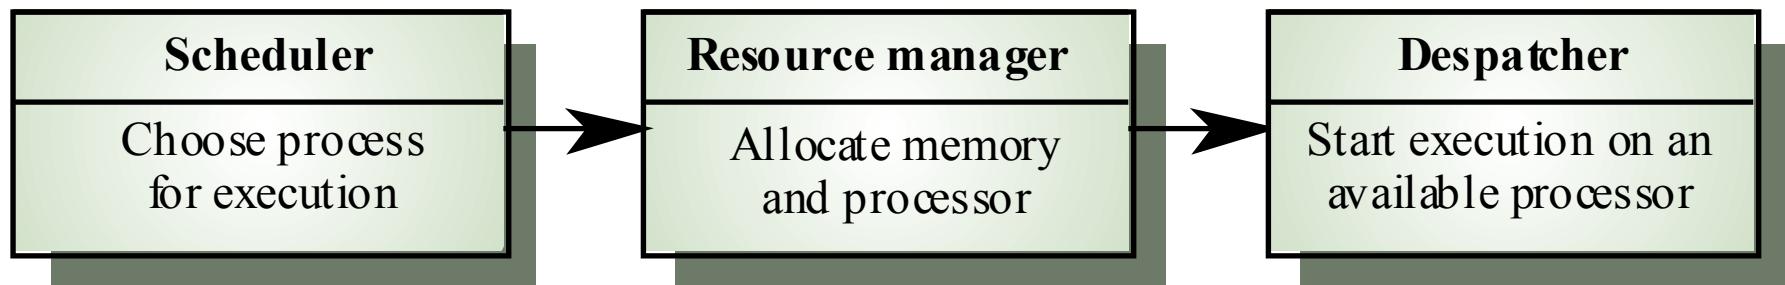


Process management

- Concerned with managing the set of concurrent processes
- Periodic processes are executed at pre-specified time intervals
- The executive uses the real-time clock to determine when to execute a process
- Process period - time between executions
- Process deadline - the time by which processing must be complete



RTE process management





Process switching

- The scheduler chooses the next process to be executed by the processor. This depends on a scheduling strategy which may take the process priority into account
- The resource manager allocates memory and a processor for the process to be executed
- The despatcher takes the process from ready list, loads it onto a processor and starts execution



Scheduling strategies

- **Non pre-emptive scheduling**
 - Once a process has been scheduled for execution, it runs to completion or until it is blocked for some reason (e.g. waiting for I/O)
- **Pre-emptive scheduling**
 - The execution of an executing processes may be stopped if a higher priority process requires service
- **Scheduling algorithms**
 - Round-robin
 - Rate monotonic
 - Shortest deadline first



Key points

- **Real-time system correctness depends not just on what the system does but also on how fast it reacts**
- **A general RT system model involves associating processes with sensors and actuators**
- **Real-time systems architectures are usually designed as a number of concurrent processes**



Key points

- **Real-time executives are responsible for process and resource management.**
- **Monitoring and control systems poll sensors and send control signal to actuators**
- **Data acquisition systems are usually organised according to a producer consumer model**
- **Java has facilities for supporting concurrency but is not suitable for the development of time-critical systems**