



Lecture 12 – Cache Memory Systems

Karl R. Wilcox
Karl@cs.rhul.ac.uk

Objectives

- In this lecture we will cover
 - Cache Memory Systems
 - The hardware for implementing cache memory

(Diagrams from Patterson & Hennessy)

Recall:

- Cache systems (of all kinds) take advantage of:
 - Temporal locality
 - If something has been referenced recently it is likely to be referenced again soon
 - Spatial locality
 - If something has been referenced, things nearby (in the address space) are likely to be referenced soon
 - Memory technology
 - Fast = expensive, slow = cheap

CPU Memory Hierarchy

- **SRAM on the same die as the CPU**
 - 5ns / Relatively large area of silicon per bit / v. high cost
- **Separate SRAM between the CPU and the bus**
 - 20ns / additional chips / high cost
- **DRAM, attached to the system bus**
 - 60ns / less silicon per bit (higher density) / medium cost
- **Magnetic disk etc. (mass storage)**
 - 1s / extremely high density / very low cost

Caching Terms

- **Objective**
 - To make as much memory as is available on the cheapest medium, at the speed of fastest
- **Cache Hit**
 - Finding an item in the cache
- **Hit time**
 - Time to access an item in the cache (including search time)
- **Cache miss**
 - Not finding an item in the cache
- **Miss penalty**
 - Time to search cache and load missing item

Compiler Considerations

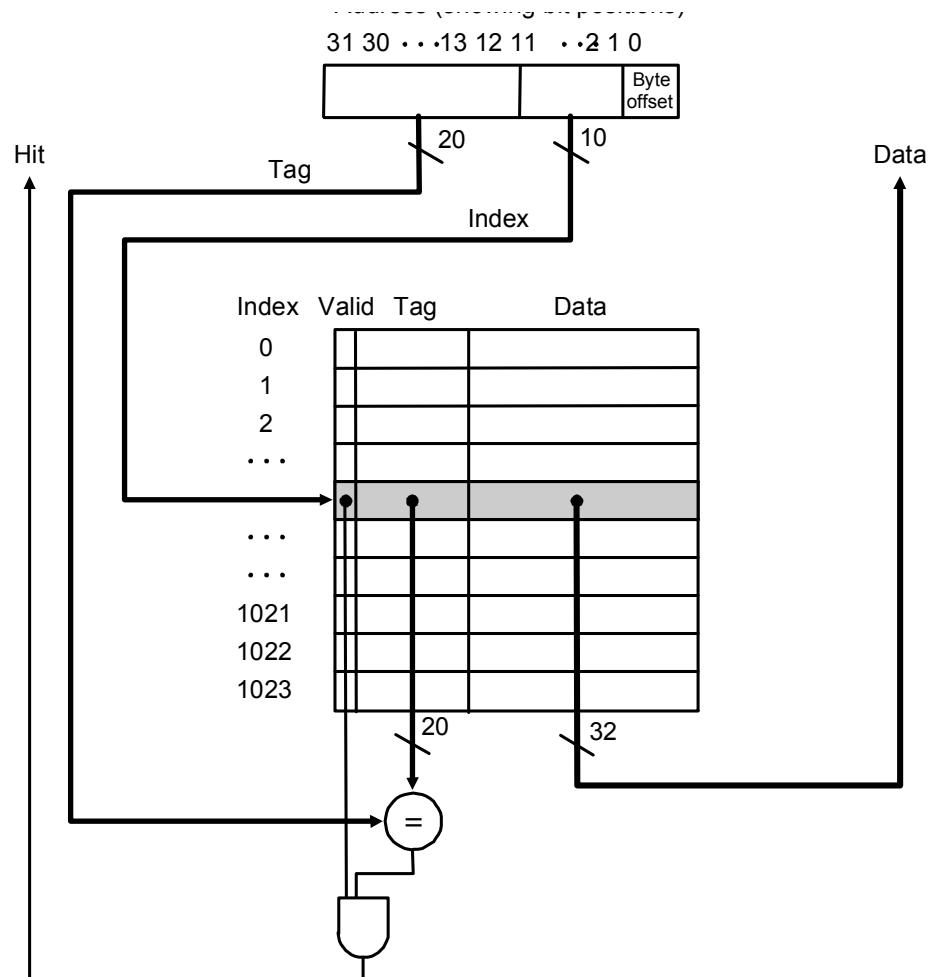
- **Although caches are in theory “invisible” to the programmer (and the compiler writer) programs need to be compiled in such a way as to maximise cache hits**
 - **I.e. need to demonstrate spatial and temporal locality**
 - **This will achieve best performance**

A Simple Memory Cache

- A cache has n blocks of storage
 - n is usually a power of 2
 - The block may be as small as a word
- *Direct Mapping*
 - Each block goes into the storage location given by the block address modulus the number of blocks in cache
 - Equivalent to the lowest m bits of the address
 - The cache must also hold the high bits to check address
 - And it requires a “valid” bit to distinguish unused locations

Hardware for a Direct Mapped Cache

- 32 bit words
- 1024, 1 word locations in cache
- Therefore
 - 2 bits for byte offset
 - 10 bits used for location
 - leaving 20 bits for tag
 - “Hit” logic needs 20 bit comparator + AND gate
- How much overhead?



The CPU and Cache Misses

- We must fetch the missing block from the next level of the memory hierarchy
- What about the CPU in the meantime?
 - If this is an instruction miss we must stall the pipeline until the instruction is available
 - If it is a data miss we could look for other instructions to execute
 - May not be worth the extra hardware...?

One Cache or Two?

- **Should we have one cache for both data and instructions?**
- **Or separate caches for each?**
- **A single cache gives a better hit rate**
 - More flexibility
- **Two caches (accessible simultaneously) increase the cache *bandwidth***
 - This usually gives better overall performance

What About Writing Data?

- **Simplest approach – write all data to memory and cache**
 - Known as *write-through*
- **Both are always valid**
- **But memory writes are slow**
- **Could write to the cache but buffer the write to memory**
 - Buffer depth can increase as required
- **Alternatively, only update cache**
 - Update memory when cache block is replaced
 - Better performance, but memory not always valid

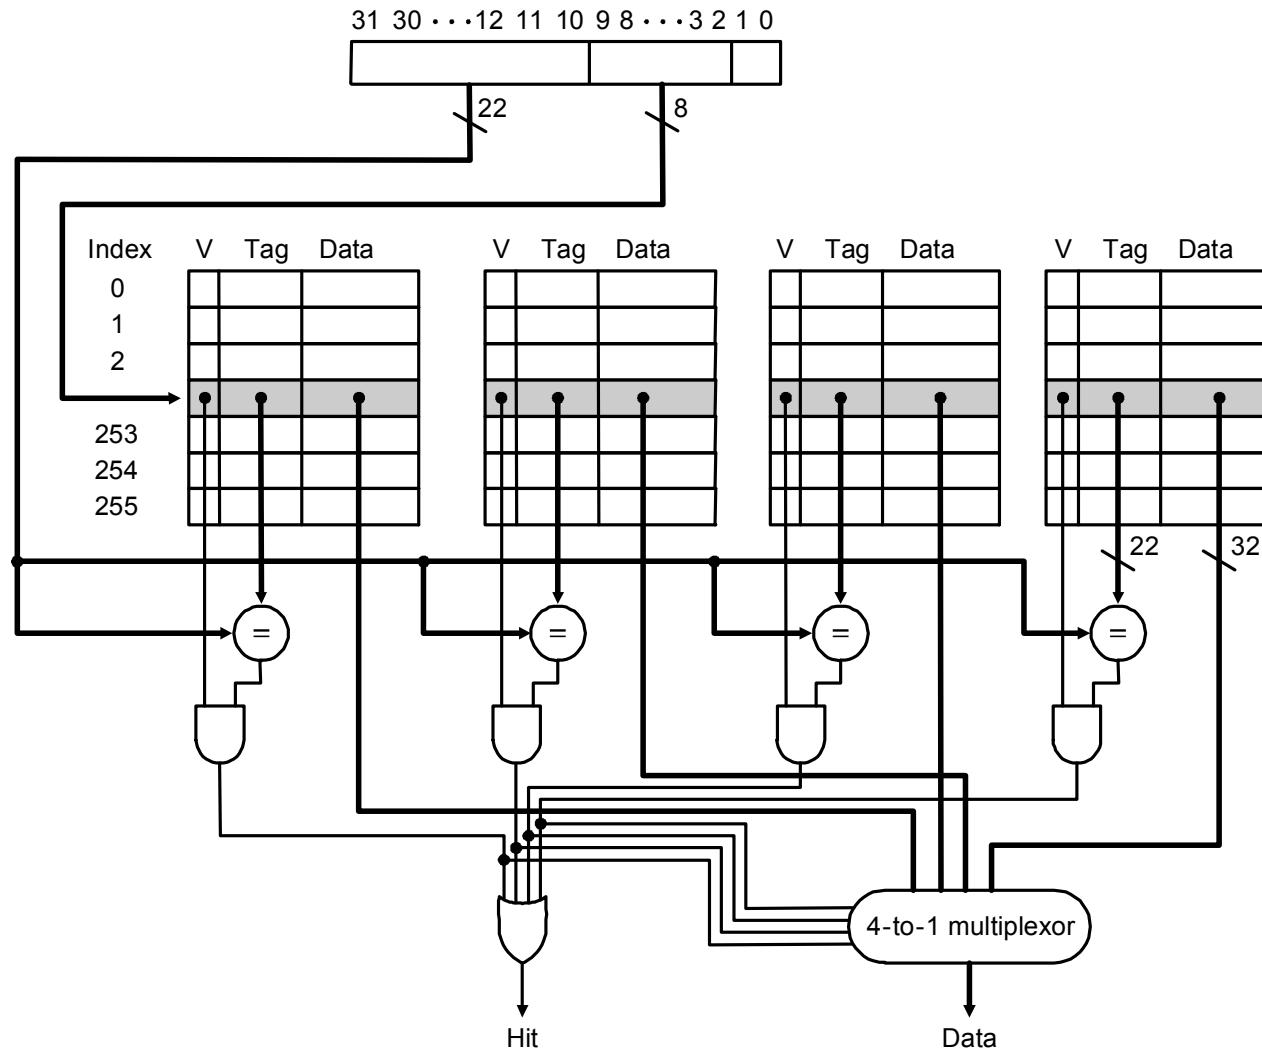
Bigger Blocks

- **One word blocks only take advantage of temporal locality**
- **Using a larger block size and loading adjacent addresses can improve performance**
- **There should be a higher hit rate**
- **But there will be a higher miss penalty**
- **We can organise the memory to read a block of data at one time**
 - Or at least more quickly – interleaving
- **As always, there is a trade-off**
 - No single answer

Other Cache Organisations

- The direct mapped cache can be wasteful
 - May leave unused locations
- Use an associative cache
 - Block can go in any location
 - Need to search all locations in parallel for target
- Compromise – the set associative cache
 - Address determines which set
 - Block can go anywhere in the set
 - Set is then searched in parallel
- Which block to replace
 - LRU is common
 - Other algorithms as discussed in CS263 Op Sys

Hardware for Set Associative Cache



Multi-level Caches

- **Modern processors often have a second level cache**
 - Separate SRAM chips “near” the processor
 - Typically 10 times as large as the primary cache
- **First level cache design should minimise hit time**
 - To give short cycle time for instructions
- **Second level cache should focus on hit rate**
 - To minimise long access times to main memory
 - Typically have bigger block sizes

Caching Vs Virtual Memory

- **Caching works over the space of a few thousand instructions**
 - I.e. “within” the quantum of time that a process gets on the processor
- **Virtual Memory operates on entire processes**
 - Which pages of a process should be in memory, ready to run
- **Caching and virtual memory are independent**
 - But both increase processor performance

Summary

- **We have looked at simple caching schemes**
- **And the hardware to implement them**
- **Some cache extensions and improvements**
 - Multi-level, set associative, write-back
- **Real caches tend to be more sophisticated**
 - But based on the same principles

Next Lecture

- **Virtual memory systems**
 - **Implementing a Virtual Memory system**
 - **(Related to CS263, Operating Systems)**
 - **Not concerned with page replacement algorithms**
 - **Interested with actual hardware implementation of virtual memory**