



Lecture 4 – Further Inter- Process Communication Techniques

Karl R. Wilcox
Karl@cs.rhul.ac.uk

Administration

- **There WILL be a lecture this Friday**
 - In MFLT, 11:00
- **Lecture Notes**
 - Are available at <http://www.cs.rhul.ac.uk/~karl>
 - Formats
 - Original Powerpoint 2000 slides
 - PDF files, A4 2 slides per page, colour
 - Let me know if you would like other formats


Objectives

- In this class we will discuss:
 - The story so far...
 - Event Channels and Monitors
 - Other Classical IPC problems

Review - 1

- **What are operating systems for?**
 - Top down view – virtual machine
 - Bottom up view – resource manager
- **Processes**
 - What is a process?
 - Process states and process pre-emption
- **Process scheduling**
 - Round Robin Scheduling
 - Priority Scheduling
 - Multiple Queues
 - Shortest Job First
 - Two-level Scheduling

Process States (Corrected)



1. Process blocks for input
2. Scheduler picks another process
3. Scheduler picks this process
4. Input becomes available

Review - 2

- **Test and Set Lock instruction**
- **Peterson's algorithm**
 - Problems with “busy waiting”
- **The producer / consumer problem & its solutions**
 - Sleep and Wakeup
 - Semaphores

Event Counters

- Event counters (a new variable type) are a tool to solve the producer-consumer problem without mutual exclusion
- Using an event counter a process can wait until a certain event has happened
- An event counter E has three operations
 - Read (E)
 - Return the current value of E
 - Advance (E)
 - Atomically increment E by 1
 - Await (E, x)
 - Wait until E has a value of x or more
- An event counter can only be incremented not decremented

P-C Solution Using Event Counters 1

```
#define N 100 // buffer size

typedef int event_counter;

event_counter in = 0;
event_counter out = 0;
```

P-C Solution Using Event Counters 2

```
void producer () {  
    int item, sequence = 0;  
  
    while (TRUE) {  
        produce_item ( &item );  
        sequence = sequence + 1;  
        await ( out, sequence - N );  
        enter_item ( item );  
        advance ( &in );  
    }  
}
```

Monitors

- A monitor is a programming language construct which obliges the programmer to declare shared data and resources explicitly and enforces mutual exclusion
- A monitor consists of
 - The data comprising a shared object
 - A set of procedures to access the object
 - Initialisation of the object (creation of object)
- The compiler (not the programmer) makes sure that mutual exclusion is guaranteed
- To ensure synchronisation the monitor provides objects called *conditions*
- Conditions have two operations
 - CWAIT (condition)
 - CSIGNAL (condition)

CWAIT and CSIGNAL

- CWAIT (condition): suspend execution of the calling process
- CSIGNAL (condition): resume execution of some process suspended after a CWAIT on the same condition
 - If there are several such processes, choose one
 - If there is no such process do nothing
- Signals to conditions are not saved (like in semaphores), they get lost
- A condition does not have a value, it can be thought of as a queue
- CWAIT adds a process to the queue
- CSIGNAL makes a process be removed from the queue if there is any

The Dining Philosophers Problem

- Five philosophers are seated around a table
- Each philosopher has a plate of spaghetti
- A philosopher needs two forks to eat it
- Between each plate is a fork
- A philosopher's life consists of alternate periods of eating and thinking
- When a philosopher gets hungry he tries to acquire his left and right fork, one at a time in either order
- if successful he eats than puts down the forks and goes on thinking
- Can you write a program for each philosopher such that each does what he is supposed to do and never gets stuck?

Why Do We Care?

- Because this is a simple example of DEADLOCK
 - If each philosopher takes the fork to his left then none of them can continue
- This is a good model of processes competing for *exclusive* access to limited resources
- Solutions to this simple problem may be solutions to real world operating systems deadlock problems

Dining Philosophers Solution - 1

```
#define N 5
#define LEFT (i-1)%N
#define RIGHT (i+1)%N
#define THINK 0
#define HUNGRY 1
#define EAT 2
typedef int semaphore;
int state[N];
semaphore mutex = 1;
semaphore s[N];
```

```
void philosopher (int i)
{
    while (TRUE)
    {
        think();
        take_forks(i);
        eat();
        put_forks(i);
    }
}
```

Dining Philosophers Solution - 2

```
void take_forks (int i)      void test(int i)
{
    down (&mutex);
    state[i] = HUNGRY;
    test(i);
    up (&mutex);
    down (&s[i]);
}

if (state[i]==HUNGRY
&& state[LEFT] !=EAT
&& state[RIGHT]!=EAT)
{
    state[i] = EAT;
    up (&s[i]);
}
```

Dining Philosophers Solution - 3

```
void put_forks(int i)
{
    down(&mutex);
    state[i]=THINK;
    test(LEFT);
    test(RIGHT);
    up(&mutex);
}
```

- **Implementation**

- All variables are shared
- mutex ensures mutual exclusion
- 5 processes
 - philosopher(1)
 - philosopher(2)
 - etc...

The Readers and Writers Problem

- Another exclusive access problem
 - Models a database system
 - Any number of simultaneous readers
 - One writer with exclusive access
- Can be solved with a semaphore for database access and a semaphore protected counter for processes reading
- Real databases normally use much finer grained locking strategies

The Sleeping Barber Problem

- **Models a “service” process with a queue of “clients”**
 - Clients are serviced in order
 - There is a fixed size queue for clients
 - If there is no room in the queue clients do not get service
 - If there are no clients, the service should sleep
- **Again, can be solved with semaphores and counters**
 - Note: our implementation of semaphores does not allow the value to be read, so we need a separate counter for clients waiting

Summary

- **Processes are fundamental to operating systems**
- **Processes are largely independent**
 - Have their own “state”
 - From programmer’s view have their own virtual machine
- **But sometimes processes interact**
 - “invisibly” as the operating system manages access to limited resources
 - “visibly” as processes need to communicate with each other
 - Producer – Consumer
 - Clients and Servers
 - Daemon Processes (see lectures on I/O)

Next Lecture

- On Friday, MFLT
- We will talk about memory management and virtual memory
- Lecture Notes: <http://www.cs.rhul.ac.uk/~karl>