



Lecture 5 – Memory Management

Karl R. Wilcox
Karl@cs.rhul.ac.uk

Administration

- **There WILL be a lecture this Friday**
 - In MFLT, 11:00
- **Lecture Notes**
 - Are available at <http://www.cs.rhul.ac.uk/~karl>
 - Formats
 - Original Powerpoint 2000 slides
 - PDF files, A4 2 slides per page, colour
 - Let me know if you would like other formats

Objectives

- In this class we will discuss:
 - Why manage memory
 - Memory allocation
 - Virtual memory
 - Paged memory
 - Page replacement algorithms

The Memory Manager

- In our discussion of processes we have assumed that all concurrent processes fit simultaneously into memory
- This is not typically the case
 - (It may be on desktop or personal systems, but still need a memory manager)
- The memory manager is responsible for:
 - Recording which parts of memory are in use or free
 - Where to allocate memory for new processes
 - How to manage swapping between main memory and disk storage

Memory Considerations

- Swapping processes out to disk may be slow
 - therefore frequent swaps are inefficient
- We aim to keep the CPU busy
 - But if in-memory processes have a lot of I/O wait time CPU utilisation may be poor
- The longer the I/O wait time, the more processes should be in memory so at least one is ready to run
 - Assume processes spend a fraction p of time in I/O wait
 - There are n processes in memory at once
 - The probability that all processes are waiting is p^n
 - CPU utilisation is $1 - p^n$

Examples

- Assume 100M of RAM, 20M for OS, user programs of 20M each
 - If each process spends 80% of its time in I/O wait:
 - CPU Utilisation = $1 - 0.8^4 = 59\%$
 - What if we had 90% I/O wait?
 - CPU Utilisation = $1 - 0.9^4 = 34\%$
 - If we add another 100M of memory (at 80% I/O)?
 - CPU Utilisation = $1 - 0.9^8 = 86\%$
 - If we add another 100M of memory (at 90% I/O)?
 - CPU Utilisation = $1 - 0.9^9 = 61\%$
 - A further 100M of memory?
 - 80% I/O = 95% 90% I/O = 77%

Memory Allocation Algorithms

- **The process of allocating main memory to processes**
 - New processes and existing processes which need additional stack or data space
- **Fixed size memory allocation**
 - May not make best use of memory
 - Easy to manage
- **Variable size memory allocation**
 - Makes better use of memory
 - More complicated to manage
- **In either case we need to track what is allocated and what is free**

Bit Maps For Memory Allocation

- **Memory is divided in *allocation units***
 - Can be a few words or several kilobytes
- **For each allocation unit there is a bit in the map**
 - The bit is 1 if the unit is occupied, 0 if it is free
- **Simple way to keep track of memory**
 - The size of the bit map depends only the size of the memory and the size of the allocation unit
- **Main problem is searching the bit map for k consecutive bits**
 - These may cross CPU word size boundaries

Linked Lists For Memory Allocation

- **Maintain a linked list of memory segments**
- **Each segment is either allocated to a process**
- **List maintained in address order**
- **Updating list is straightforward**
- **Adjacent holes must be merged (to combine them into a single larger hole)**
- **Therefore it is easier if the list is doubly linked**

Link List Algorithms

- **It is easy to find all available holes in the list**
 - Which of them should be allocated?
- **First Fit**
 - Scan the list until one is found sufficiently large
- **Next Fit**
 - As first fit, but starting from where you left off
- **Best Fit**
 - Allocate smallest hole that the process fits in
- **Worst Fit**
 - Allocate largest available hole
- **Quick Fit**
 - Maintain separate lists for common sizes

Buddy System For Memory Allocation

- The buddy system takes advantage of computers using binary numbers for addressing
- The manager maintains a list of free blocks of size 1, 2, 4, 8, 16, etc. bytes up to size of memory
- Initially, one hole the size of the whole memory
- To allocate space for a process which needs x bytes, find a hole of size 2^k , where $2^{k-1} < x \leq 2^k$
- If necessary bigger holes are divided into two holes half as big

Not Enough Memory

- The previous discussion assumed that processes would fit into the available memory
 - Not always the case
- An early solution – Overlays
 - Overlay 0 runs first
 - When done calls next overlay into memory
 - Overlay splitting was done by the programmer
 - Performance depended on how well (how modular) the split was
- Replaced by the concept of virtual memory

Virtual Memory

- The combined size of the program, data, and stack may exceed the amount of physical memory available
- The operating system keeps those parts of the program currently in use in the main memory
- The rest is kept on disk
- Applicable to all the processes in a multi-programming system

Paging

- Virtual memory systems use the technique Paging
- On a computer there exists a set of memory addresses that programs can produce
- They are called virtual addresses
- They form the virtual address space
- These addresses go to a memory management unit (MMU)
- The MMU hardware maps the virtual addresses onto the physical memory address

Pages and Page Frames

- The virtual space is divided into fixed sized pages
- The physical memory is divided into page frames of the same size
- Transfers between the two memories are always in units of a page

Physical
Memory



Disk
Memory



Page Faults

- Since the virtual memory is bigger than the physical memory not every page is mapped onto a page frame
- A present / absent bit in each entry keeps track of whether a page is mapped or not
- When the program tries to use an unmapped page it causes a trap to the operating system
- This trap is called a page fault
- The OS chooses one page frame, writes the contents back to disk and fetches the called page into the page frame, then updates the MMU map

Page Tables

- **Page Tables** are used for the mapping (virtual address → physical address) in the MMU
- The virtual address is split into a virtual page number and offset
 - Virtual page number = index into page table
 - Offset = last part of (virtual and physical) address
- The major issues:
 - The page table can be extremely large
 - The mapping must be fast

Multilevel Page Tables

- To avoid keeping all page tables in memory
 - Uses a top level page table
 - And second level page tables
- The virtual address is now partitioned into:
 - PT1 field (e.g. 10 bits)
 - PT2 field (e.g. 10 bits)
 - Offset field (e.g. 12 bits)
- The PT1 field indexes the top level page table
- There we get the address of the second level page table (indexed by the PT2 field)
- This table contains the physical address

Associative Memory

- Most programs tend to make a large number of references to a small number of pages
- Therefore computers are equipped with a small hardware device for mapping virtual addresses to physical addresses without using page tables
- This device is called an associative memory
- Has a small number of entries, each about 1 page
- The OS looks up the page in the associative memory first
- If not present then the normal page tables are consulted

Summary

- **We wish to have multiple processes in memory**
 - So that at least one is ready to run
- **Therefore memory must be shared between processes**
 - Need to keep track of which memory is used by what
- **We would also like to let processes run that need more memory than physically available**
 - We use virtual paged memory to achieve this

Next Lecture

- On Friday, MFLT
- We will talk about page replacement algorithms
- Lecture Notes: <http://www.cs.rhul.ac.uk/~karl>